

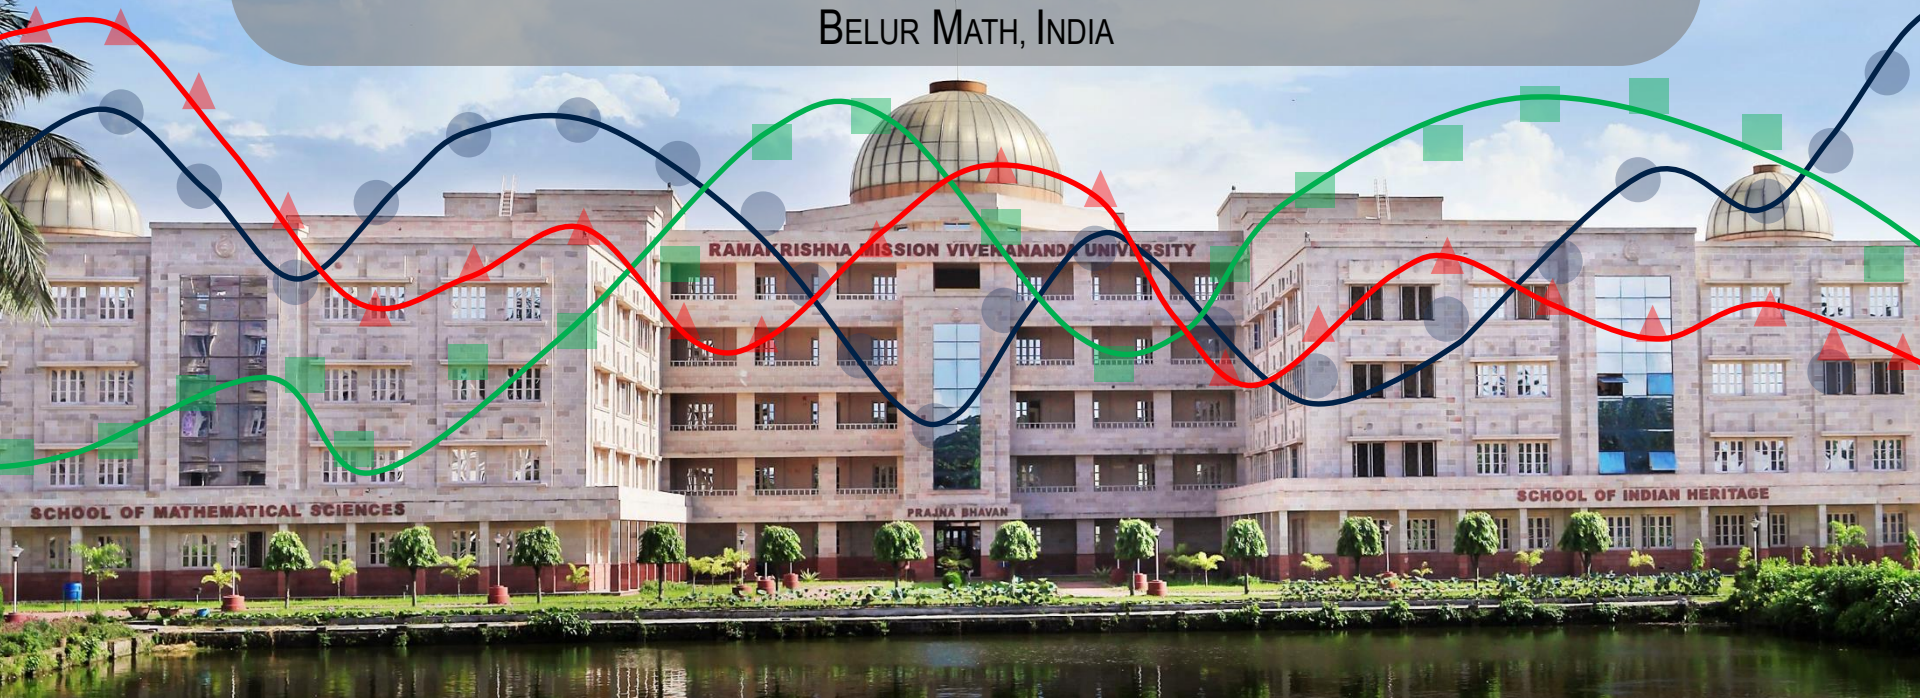
Machine Learning: The Basics

DRIPTA MJ

Department of Mathematics

RAMAKRISHNA MISSION VIVEKANANDA EDUCATIONAL AND RESEARCH INSTITUTE

BELUR MATH, INDIA



Machine Learning is everywhere

Astronomy



Social Networks



Healthcare



Banking




Genomics




Weather predictions



Dogs and Cats




?




?



Product recommendation



+




+



Images from *amazon.com*

Autonomous vehicles





Creativity





Figure source: Gatys, Ecker and Bethge, Image style transfer using convolutional neural networks, CVPR 2016.

Machine Learning

SUPERVISED



UNSUPERVISED




Some key components

Data pre-processing

x_1	x_2	x_3	y
2.2	0.8	2.7	1
4.9	3.1	1.6	-1

- Data cleaning
- Training-test data splitting
- Feature engineering

ML Model



Linear Regression

k NN

SVM

Decision Tree

Neural Network


.....

Training



- Loss function
- Optimization algorithm
- Regularization

Evaluation



- Generalization error
- Cross-validation
- Metric

Features

- Attributes used to represent input data.
- Features of *Iris* species:
 - Sepal Length
 - Sepal Width
 - Petal Length
 - Petal Width



Iris dataset


INPUTS

Sepal Length (cm)	Sepal Width (cm)	Petal Length (cm)	Petal Width (cm)
5.1	3.5	1.4	0.2
4.9	3	1.4	0.2
4.7	3.2	1.3	0.2
4.6	3.1	1.5	0.2
5	3.6	1.4	0.2
5.4	3.9	1.7	0.4
4.6	3.4	1.4	0.3
5	3.4	1.5	0.2
4.4	2.9	1.4	0.2
.	.	.	.
.	.	.	.

OUTPUTS

Species	
Iris Setosa	0
Iris Virginica	1
Iris Versicolor	2


Training and Test data



- Training data: Used for training the ML algorithm.
- Test data: Used for assessing the performance of the ML algorithm.

Loss function


REGRESSION



Squared loss:

$$\mathcal{L}(\mathbf{y}^{(n)}, \mathbf{y}^{*(n)}) = \frac{1}{2} \sum_{j=1}^J (y_j^{(n)} - y_j^{*(n)})^2$$


BINARY CLASSIFICATION



Binary cross-entropy loss:

$$\mathcal{L}(y^{(n)}, y^{*(n)}) = -y^{(n)} \log(y^{*(n)}) - (1 - y^{(n)}) \log(1 - y^{*(n)})$$


MULTI-CLASS CLASSIFICATION



Cross-entropy loss:

$$\mathcal{L}(\mathbf{y}^{(n)}, \mathbf{y}^{*(n)}) = - \sum_{j=1}^J y_j^{(n)} \log y_j^{*(n)}$$


Generalization




- Larger class of functions \rightarrow more complexity of the hypothesis class $\mathcal{C}(\mathbb{H})$.
- Objective: Good prediction at unobserved locations \rightarrow good **generalization**.

Generalization

REGRESSION




CLASSIFICATION




Figures for illustration only.

Simple models

REGRESSION




CLASSIFICATION




Figures for illustration only.

Complex models

REGRESSION




CLASSIFICATION



Figures for illustration only.

Loss vs complexity



Figures for illustration only.

Bias-variance decomposition

- Dataset: $\mathcal{D} = \left\{ (\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)}) \right\}$
- Let $g_{\mathcal{D}}$ be the hypothesis which is fit to a particular training dataset \mathcal{D}
- Want to compute the expected prediction error at an arbitrary test point with input \mathbf{x} and output y : $\mathbb{E}_{\mathbf{x}, y, \mathcal{D}} \left[(g_{\mathcal{D}}(\mathbf{x}) - y)^2 \right]$.

- Mean prediction of the machine learning algorithm:

$$\bar{g}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}} \left[g_{\mathcal{D}}(\mathbf{x}) \right]$$

- So determining the value of $\bar{g}(\mathbf{x})$ involve
 - generating different training datasets (\mathcal{D}),
 - training separate functions ($g_{\mathcal{D}}$) for every generated dataset,
 - making predictions at an arbitrary test point \mathbf{x} with all trained functions,
 - and finally, averaging over all the predictions.
- Let $\bar{y}(\mathbf{x})$ be the expected value of the output at \mathbf{x} , i.e. $\bar{y}(\mathbf{x}) = \mathbb{E}_{y|\mathbf{x}}[y]$.

Bias-variance decomposition

- Dataset: $\mathcal{D} = \left\{ (\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)}) \right\}$
- Let $g_{\mathcal{D}}$ be the hypothesis which is fit to a particular training dataset \mathcal{D}
- Want to compute the expected prediction error at an arbitrary test point with input \mathbf{x} and output y : $\mathbb{E}_{\mathbf{x}, y, \mathcal{D}} \left[(g_{\mathcal{D}}(\mathbf{x}) - y)^2 \right]$.

- Mean prediction of the machine learning algorithm:

$$\bar{g}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}} \left[g_{\mathcal{D}}(\mathbf{x}) \right]$$

- So determining the value of $\bar{g}(\mathbf{x})$ involve
 - generating different training datasets (\mathcal{D}),
 - training separate functions ($g_{\mathcal{D}}$) for every generated dataset,
 - making predictions at an arbitrary test point \mathbf{x} with all trained functions,
 - and finally, averaging over all the predictions.
- Let $\bar{y}(\mathbf{x})$ be the expected value of the output at \mathbf{x} , i.e. $\bar{y}(\mathbf{x}) = \mathbb{E}_{y|\mathbf{x}}[y]$.


Bias-variance decomposition

$$\mathbb{E}_{\mathbf{x}, y, \mathcal{D}} \left[(g_{\mathcal{D}}(\mathbf{x}) - y)^2 \right] = \underbrace{\mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(g_{\mathcal{D}}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2 \right]}_{\text{Variance}} + \underbrace{\mathbb{E}_{\mathbf{x}} \left[(\bar{g}(\mathbf{x}) - \bar{y}(\mathbf{x}))^2 \right]}_{\text{Bias}^2} + \underbrace{\mathbb{E}_{\mathbf{x}, y} \left[(\bar{y}(\mathbf{x}) - y)^2 \right]}_{\text{Noise}}$$

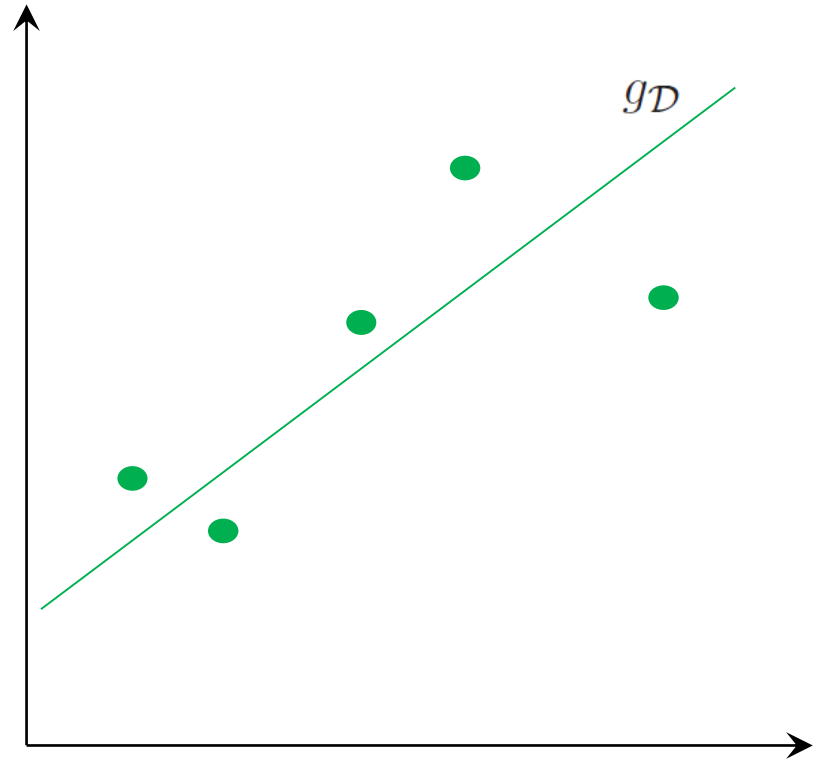
- **Variance**: It expresses the sensitivity of the solution on the particular choice of dataset \mathcal{D} .
- **Bias**: Difference between the expected prediction (averaged over different datasets) and the expected output value. This is the inherent error arising from the choice of model.
- **Noise**: Expresses the noise in the data.

Example

Underlying true function f
and data points






Hypothesis fit: $g_{\mathcal{D}}$



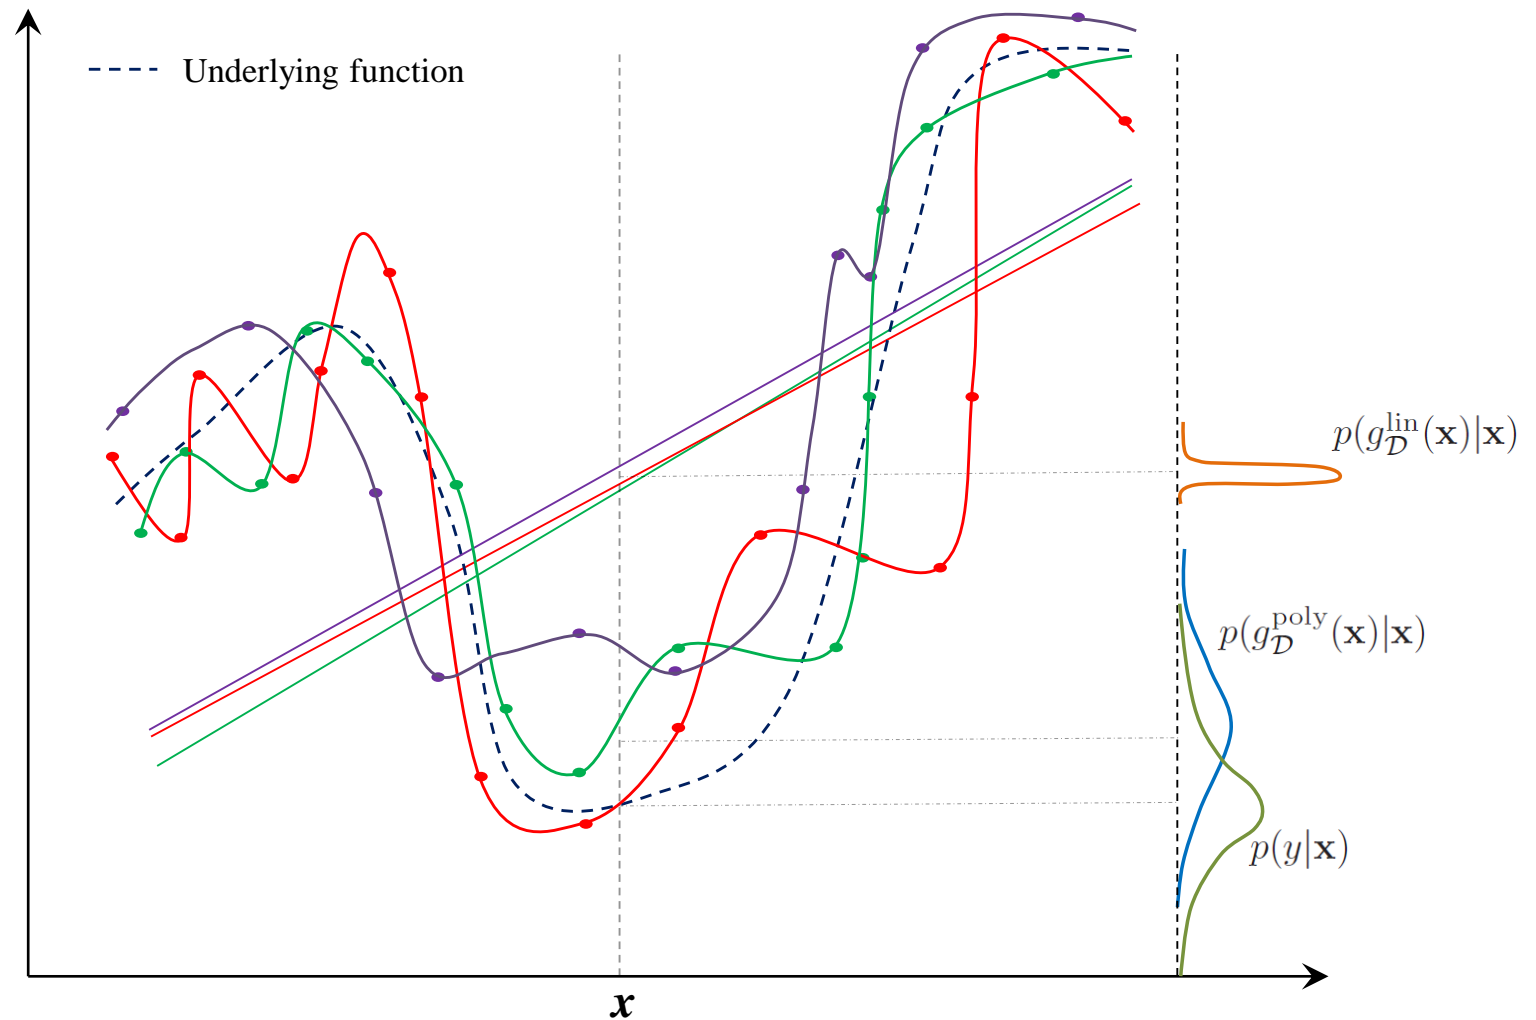
Figures for illustration only.

Visualization




Figures for illustration only.

Another example



Figures for illustration only.


Bias, variance vs model complexity



- **High Bias:** Model is too simple, and so unable to fit the data properly.
 - Results in underfitting.
 - Training and test errors are both large.
- **High Variance:** Model is too complex, and so small changes in the data produce significant changes in the solution.
 - Results in overfitting.
 - Test Error \gg Training Error

Figures for illustration only.

Underfitting & Overfitting



- **Underfitting** can be addressed by
 - Increasing the complexity of the model.
 - Minimizing the cost function properly in the training stage.
- **Overfitting** can be addressed by
 - Reducing the complexity of the model.
 - Incorporating some form of regularization inside the cost function.

Figures for illustration only.

Training and Test datasets


- Dataset is split into two groups:
 - Training dataset is used to train the ML algorithm.
 - Test dataset is used to estimate the error rate of the trained model.



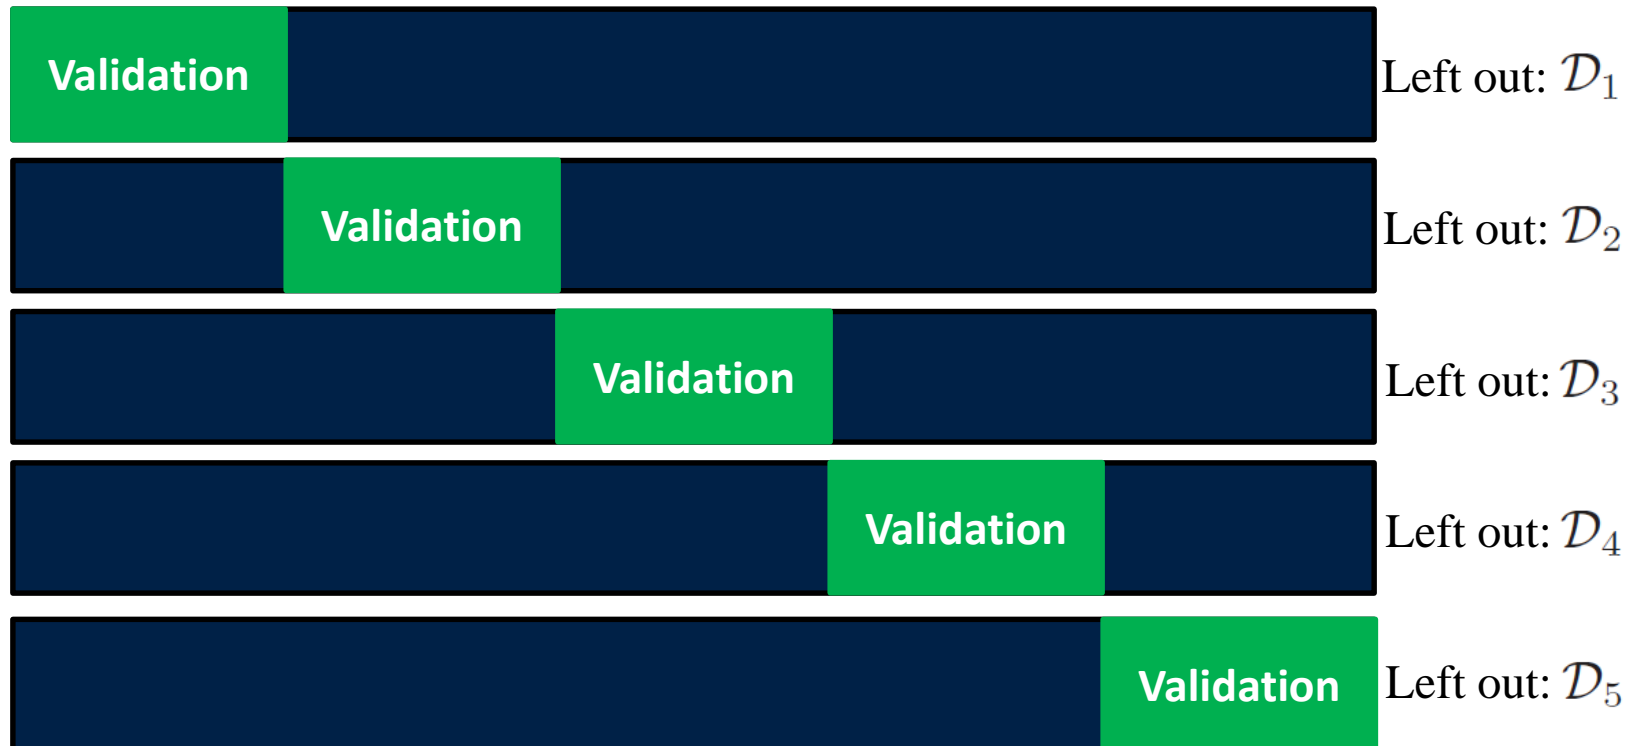
- Shortcomings:
 - If the size of the dataset is small, then keeping aside a separate test dataset can lead to loss of some vital information in the model training stage.
 - “Unfortunate” data split can result in misleading error estimates.
- Solution:
 - K -fold cross-validation
 - Leave-one-out cross-validation

K-fold cross validation

- Training data is subdivided into K separate subsets – $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_K$ of equal size (say n_K). Let's take $K = 5$.



- Can generate K training-test datasets using the K subsets



K-fold cross validation

- For $k = 1, 2, \dots, K$
 - Leave out the k th fold data \mathcal{D}_k and train the model on the remaining $k - 1$ folds.



- Use the trained model to make prediction on the k th fold data \mathcal{D}_k and compute the (cross validation) error for this fold

$$E_k = \frac{1}{n_K} \sum_{i=1}^{n_K} (y_{k,i} - f_{-k}(\mathbf{x}_i))^2$$

where f_{-k} is the model trained excluding the k th fold data \mathcal{D}_k .


K-fold cross validation

- Estimated generalization error:

$$\mathbf{E} = \frac{1}{K} \sum_{k=1}^K E_k$$


- When $K = N$ (size of the training dataset), the approach is known as [leave-one-out cross-validation](#).
- **Note:** Cross-validation is also used to [tune the hyperparameters](#) of a model.
 - The optimal value of a hyperparameter is the one yielding the least value of \mathbf{E} .

Optimization




- Objective: $\min_{\theta} \mathcal{L}(\theta)$

Contours and gradients



- Contour line is a level curve which is the set of all real-valued solutions for a fixed value of the objective function.
- Gradients are perpendicular to contour lines.

Gradient Descent algorithm




- Gradient descent algorithm (1st order method)

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} - \xi_k \nabla g(\boldsymbol{\theta}^{(k)})$$


where k is the iteration no. and $\xi_k > 0$ is the learning rate or step size.

- Note that the gradients point towards the maximum, and therefore a negative sign is introduced in front of ξ_k since we are interested in the minimum.

Step size




Fixed step-size: ξ_a



Fixed step-size: ξ_b

$$\xi_a < \xi_b$$

Regularization



Figures for illustration only.

Classification: Confusion matrix

- A table used to describe/visualize the performance of a classification algorithm.
- Confusion matrix for a binary classification problem:

		Prediction	
		Negative	Positive
Actual	Negative	980	6
	Positive	4	10

- Standard metric: Accuracy
 - Ratio of number of correct predictions to all predictions.

$$\text{Accuracy} = \frac{\text{True Positive} + \text{True Negative}}{\text{True Positive} + \text{False Positive} + \text{True Negative} + \text{False Negative}}$$